

PROJECT REPORT ON "Predictive Modeling of

✓ Cardiovascular Disease: A Comparative Analysis of Machine Learning Algorithms"

1. Project Objective

- 1) Identification of High-Performance Classification Models: The primary objective of this project is to evaluate the effectiveness of supervised algorithms classification in predicting cardiovascular diseases using the provided dataset.
- 2) Comparative Analysis with Alternative Models: Another goal is to compare the performance of the decision tree classifier with other popular classification algorithms, namely K-Nearest Neighbors (KNN) and Support Vector Machines (SVM). By employing these models, we aim to determine which algorithm yields the highest predictive accuracy for cardiovascular disease prediction.
- 3) Identification of Key Predictive Features: Beyond model performance, this project aims to identify the most influential features in predicting cardiovascular diseases. Through feature importance analysis and threshold determination, we seek to ascertain the variables that significantly contribute to accurate classification.

2. Description of Data

2.1 Data Source, Size, Shape

2.1.1 Data Source (Website Link): The dataset was obtained from Kaggle, a widely recognized platform for data science resources and competitions. The dataset can be accessed at the following link: Cardiovascular Disease Dataset on Kaggle

***2.1.2 Data Size:** *The dataset size is appropriate for analysis, occupying a moderate storage space.

2.1.3 Data Shape (Dimension: Number of Variables | Number of Records): The dataset comprises 70,000 entries and consists of 13 columns, providing a comprehensive set of health-related information for each individual.

2.2 Variable Description

2.2.1 Index Variable: The 'id' column functions as the index variable, offering a unique identifier for each individual in the dataset. This variable is excluded from the analytical process as it does not contribute to health-related insights.

2.2.2 Categorical Variables:

Nominal Type: 'gender', 'cholesterol', 'gluc', 'smoke', 'alco', 'active'. Ordinal Type: None.

2.2.3 Non-Categorical Variables:

'age': Age of the individual in days. 'height': Height of the individual in centimeters. 'weight': Weight of the individual in kilograms. 'ap_hi': Systolic blood pressure. 'ap_lo': Diastolic blood pressure.

3.1 # Source of data

The dataset was sourced from Kaggle, a renowned platform renowned for hosting data science competitions, datasets, and tutorials

[link text](#)

Variable Description:

3.2 Index Variable:

The 'id' column serves as the index variable, providing a unique identifier for each individual. This variable will not be included in the analysis section as it does not contribute to the health-related insights.

3.3 Categorical Variables:

Nominal Type: 'gender', 'cholesterol', 'gluc', 'smoke', 'alco', 'active'. Ordinal Type: None.

3.4 Non-Categorical Variables:

'age': Age of the individual in days. 'height': Height of the individual in centimeters. 'weight': Weight of the individual in kilograms. 'ap_hi': Systolic blood pressure. 'ap_lo': Diastolic blood pressure.

3.5 Numerical description of non categorical variables-

interpretation of each variable's mean, median, and mode:

3.5.1 Age:

Mean: The average age of individuals in the dataset is approximately 19468.87 days.

Median: The middle value of ages when arranged in ascending order is about 19703 days.

Mode: The most frequent age in the dataset is 18236 days, indicating it's the most common age among individuals.

3.5.2 Height:

Mean: The average height of individuals in the dataset is around 164.36 cm.

Median: The middle value of heights when arranged in ascending order is 165 cm.

Mode: The most frequent height observed is 165 cm, suggesting it's the most common height among individuals.

3.5.3 Weight:

Mean: The average weight of individuals in the dataset is approximately 74.21 kg.

Median: The middle value of weights when arranged in ascending order is 72 kg.

Mode: The most frequent weight observed is 65 kg, indicating it's the most common weight among individuals.

3.5.4 Systolic Blood Pressure (ap_hi):

Mean: The average systolic blood pressure is approximately 128.82 mmHg.

Median: The middle value of systolic blood pressure when arranged in ascending order is 120 mmHg.

Mode: The most frequent systolic blood pressure observed is 120 mmHg, suggesting it's the most common value among individuals.

3.5.5 Diastolic Blood Pressure (ap_lo):

Mean: The average diastolic blood pressure is approximately 96.63 mmHg.

Median: The middle value of diastolic blood pressure when arranged in ascending order is 80 mmHg.

Mode: The most frequent diastolic blood pressure observed is 80 mmHg, indicating it's the most common value among individuals.

3.6 Frequency info and relative freq. info for categorical variables.

3.6.1 Gender:

Count:

Male: 45530 Female: 24470 Percentage:

Male: 65% Female: 35%

Interpretation: There are 45530 (65%) males and 24470 (35%) females in the dataset.

3.6.2 Cholesterol:

Count:

Normal: 52385 Above Normal: 9549 Well Above Normal: 8066

Percentage: Normal: 75% Above Normal: 14% Well Above Normal: 12%

Interpretation: The majority of individuals (75%) have normal cholesterol levels, while a smaller proportion have above normal (14%) and well above normal (12%) cholesterol levels.

3.6.3 Glucose Level (gluc):

Count:

Normal: 59479 Above Normal: 5190 Well Above Normal: 5331

Percentage:

Normal: 85% Above Normal: 7% Well Above Normal: 8%

Interpretation: Most individuals (85%) have normal glucose levels, with smaller proportions having above normal (7%) and well above normal (8%) glucose levels.

3.6.4 Smoking Habit (smoke):

Count:

Non-Smoker: 63831 Smoker: 6169

Percentage: Non-Smoker: 91% Smoker: 9%

Interpretation: The majority of individuals (91%) are non-smokers, while a smaller proportion (9%) are smokers.

3.6.5 Presence of Cardiovascular Disease (cardio):

Count:

Absent: 35021 Present: 34979

Percentage:

Absent: 50% Present: 50%

Interpretation: The dataset is evenly split between individuals with and without cardiovascular disease, each comprising 50% of the dataset.

3.6.6 Physical Activity Level (active):

Count:

Active: 56261 Inactive: 13739

Percentage: Active: 80% Inactive: 20%

Interpretation: The majority of individuals (80%) are physically active, while a minority (20%) are inactive.

3.6.7 Alcohol Consumption (alco):

Count:

Non-Drinker: 66236 Drinker: 3764

Percentage: Non-Drinker: 95% Drinker: 5%

Interpretation: The majority of individuals (95%) do not consume alcohol, while a small proportion (5%) are drinkers.

4. **Data Preprocessing:

4.1 Missing Data Info & Treatment:**

4.1.1 Missing Data Info:

There are no missing values in either the categorical or non-categorical variables.

4.1.2 Missing Data Treatment:

No,records have missing data. No,variables have missing data.

4.2 Categorical Data Numerical Encoding:

Alpha-Numeric Variables:All categorical variables are already encoded numerically.

Alpha-Numeric Numerical Assignment:No alpha-numeric numerical assignment is needed as the data is already encoded numerically.

4.3 Non-Categorical Data Outliers Detection and Treatment:

Outliers Detection:Box plots were generated for non-categorical variables ('age', 'height', 'weight', 'ap_hi', 'ap_lo') to detect outliers. Outliers Treatment:

Outliers detected using box plots were treated by normalization using Min-Max scaling.

4.4 Bifurcation of Data into Training and Testing:

Bifurcation:Bifurcation into training and testing datasets was not performed, as it's not applicable for clustering analysis.

Selection of Variables:Only numerical encoded variables were considered for further analysis.

5. Supervised Machine Learning Classification Algorithm: Decision Tree (Base Model)

5.1 Decision Tree Model (Base Model) - Metrics Used: Gini Coefficient, Entropy

*The decision tree model was trained using the provided cardiovascular dataset with the aim of segmenting individuals into different clusters based on their health-related attributes. Two metrics, namely the Gini coefficient and entropy, were utilized to evaluate the performance of the decision tree model.

*The decision tree model, with a maximum depth of 3, was trained on the training subset of the dataset. The resulting model rules were extracted, providing a clear insight into the hierarchical decision-making process employed by the decision tree classifier. The rules indicated specific thresholds for different features such as systolic blood pressure, cholesterol levels, and height, among others, which contributed to the classification of individuals into distinct clusters.

- Decision Tree Model Rules:

```
|-- cardio_y <= 0.50 | |-- ap_hi_mn <= 0.02 | | |-- ap_hi_mn <= 0.02 | | | |-- class: 1 | | |-- ap_hi_mn > 0.02 | | | |-- class: 0 | |-- ap_hi_mn > 0.02 | | |-- cholesterol_y <= 1.50 | | | |-- class: 0 | | | |-- cholesterol_y > 1.50 | | | |-- class: 2 |-- cardio_y > 0.50 | |-- cholesterol_y <= 1.50 | | |-- height_mn
```

```
<= 0.53 ||| |-- class: 2 ||| |-- height_mn > 0.53 ||| |-- class: 0 ||| |-- cholesterol_y > 1.50 ||| |--  
height_mn <= 0.53 ||| |-- class: 2 ||| |-- height_mn > 0.53 ||| |-- class: 2
```

- These rules provide a clear understanding of the decision boundaries created by the decision tree model, enabling the classification of individuals into different clusters based on their health attributes.

Feature Importance:

Cardio_y: This feature has the highest importance score of 0.54, indicating that it plays a significant role in predicting the outcome. It likely represents some cardiovascular health metric or condition.

Ap_hi_mn: This feature comes second in importance with a score of 0.33. "Ap_hi_mn" could refer to average high blood pressure (systolic blood pressure) over time, and its relatively high importance suggests that blood pressure might be a crucial factor in predicting the outcome.

Cholesterol_y: While less important than the previous two features, cholesterol levels (cholesterol_y) still contribute significantly to the predictive power of the model with an importance score of 0.09. High cholesterol is a known risk factor for cardiovascular diseases.

Height_mn: Height has a relatively lower importance score compared to the previous features (0.04), suggesting it contributes less to predicting the outcome in this context.

Gender_y, Gluc_y, Smoke_y, Alco_y, Active_y, Weight_mn, Ap_lo_mn: These features all have importance scores of 0.00, indicating that they are not contributing much, if at all, to predicting the outcome in this model. It doesn't mean they are irrelevant overall, but rather in this specific predictive model, they don't play a significant role.

5.1 Supervised Machine Learning Classification Algorithms: Comparison Models

The trained decision tree model was then compared with these alternative models using metrics such as precision, recall, and F1-score. The F1-score, which is the harmonic mean of precision and recall, provides a balanced measure of a model's performance, especially in situations where there is an imbalance between the classes.

5.1.1 Decision Tree Model Performance (Testing Subset):

The confusion matrix and classification report were generated to evaluate the decision tree model's performance on the testing subset of the dataset. The confusion matrix provides a tabular representation of the true positive, true negative, false positive, and false negative predictions made by the model. Additionally, the classification report presents metrics such as precision, recall, F1-score, and support for each class, providing a comprehensive assessment of the model's predictive capabilities.

5.1.2 Interpretation of F1-score:

The F1-score is a measure of a model's accuracy that considers both the precision and recall of the model. The F1-score indicates the harmonic mean of precision and recall for each class in the classification task. A higher F1-score indicates better model performance, with values closer to 1 indicating a perfect balance between precision and recall.

The F1-score is the harmonic mean of precision and recall and provides a balance between the two metrics. In this report:

For class 0, the F1-score is 0.53, which combines precision and recall for class 0. For class 1, the F1-score is 0.32, representing the balance between precision and recall for class 1. For class 2, the F1-score is 0.35, indicating the balance between precision and recall for class 2.

The model achieves an overall accuracy of 44%. Class 0 has the highest support, followed by Class 2 and Class 1, indicating an imbalanced class distribution.

For Class 0, precision is 42%, recall is 71%, and F1-score is 0.53, showing a balanced performance.

For Class 1, precision is 44%, recall is 25%, and F1-score is 0.32, indicating lower recall and imbalance.

For Class 2, precision is 52%, recall is 26%, and F1-score is 0.35, suggesting moderate balance but relatively low overall performance.

The weighted average metrics are higher than macro-average, indicating the influence of class distribution, with Class 0 having a greater impact.

5.2 Confusion Matrix

A confusion matrix is a table that visually represents the performance of a classification model. It compares the actual labels for data points with the labels predicted by the model. In the image, rows represent the actual labels, and columns represent the predicted labels. The value at each cell represents the number of data points that belong to a particular class (shown in the rows) and were predicted to belong to the class indicated by the column. Ideally, most of the values should be concentrated on the diagonal cells, which means the model correctly classified the data points.

5.2.1 Interpreting the Confusion Matrix

Class 0: The model performed well on class 0, correctly classifying 3,000 out of 3,151 data points. There were 349 misclassifications, with most being classified as class 1 (290). Class 1: The model's performance on class 1 was moderate. Out of 2,394 data points in class 1, the model correctly classified 1,200. It misclassified 873 as class 0 and 321 as class 2. Class 2: The model performed poorly on class 2. Out of 4057 data points in class 2, the model only correctly classified 400. It incorrectly classified a significant portion (3,500) as class 0 and 157 as class 1.

Key Takeaways

The decision tree classifier seems to be biased towards class 0, with a high number of correct classifications for this class. The model struggles to distinguish between class 1 and class 2, misclassifying a substantial number of data points from these classes.

5.2.2 Conclusion:

Through the analysis of the decision tree model and comparison with alternative classification algorithms, valuable insights were gained into the segmentation of individuals based on their health attributes. The decision tree model, with its interpretable rules and competitive performance metrics, emerges as a promising approach for segmenting individuals into clusters based on cardiovascular attributes.

This analysis provides a comprehensive understanding of the segmentation of individuals based on their cardiovascular attributes using machine learning classification algorithms, as outlined in the provided code snippets and outputs.

6. Support Vector Machine

The Support Vector Machine (SVM) classifier was employed to segment individuals into clusters based on their health-related attributes. The SVM model was trained using a linear kernel on the training subset of the dataset and evaluated using the testing subset.

6.1 SVM Model Performance (Testing Subset):

The confusion matrix and classification report were generated to assess the SVM model's performance. The confusion matrix illustrates the true positive, true negative, false positive, and false negative predictions made by the model. Additionally, the classification report provides metrics such as precision, recall, F1-score, and support for each class.

6.2 Interpretation of F1-score:

The F1-score is a measure of a model's accuracy that considers both precision and recall. A higher F1-score indicates better model performance, with values closer to 1 indicating a better balance between precision and recall.

The F1-score is the harmonic mean of precision and recall and provides a balance between the two metrics. In this report:

For class 0, the F1-score is 0.47, which combines precision and recall for class 0. For class 1, the F1-score is 0.00, representing the balance between precision and recall for class 1. Since there were no true positive predictions for class 1, the F1-score is also 0. For class 2, the F1-score is 0.51, indicating the balance between precision and recall for class 2.

6.3 SVM Model Evaluation Results:

The SVM model achieved an overall accuracy of 43% on the testing subset. However, it's important to note that the precision, recall, and F1-score varied across different classes. For example, the F1-score for class 1 was reported as 0.00, indicating that the model had difficulty

correctly classifying individuals belonging to this class. On the other hand, the F1-score for class 2 was 0.51, suggesting relatively better performance in identifying individuals in this category.

**6.4 Decision Tree:* *This portion of the image visualizes the decision-making process of the SVM classifier. It reveals a series of binary questions asked about the data to arrive at a classification. Each branch in the tree represents a possible answer to the question, and the terminal nodes represent the final classifications.

6,5 Confusion Matrix: The confusion matrix is a table that allows visualization of the performance of the SVM. It compares the actual labels for the data points with the labels predicted by the model. In the image, rows represent the actual labels, and columns represent the predicted labels. The value at each cell represents the number of data points that belong to a particular class (shown in the rows) and were predicted to belong to the class indicated by the column. Ideally, most of the values should be concentrated on the diagonal cells, which means the model correctly classified the data points.

The model seems to have performed well on class 0, correctly classifying 3,000 out of 3,000 data points. The model's performance on class 1 was not as strong. Out of 2,700 data points in class 1, the model only correctly classified 1,200. It incorrectly classified 1,500 as class 2. The model performed poorly on class 2. None of the 2,300 data points in class 2 were correctly classified. The model incorrectly classified all of them as class 0. Overall, the confusion matrix suggests that the SVM model might be biased towards class 0 and needs some improvement in classifying class 1 and 2.

7. K-Nearest Neighbors (KNN) Model

The K-Nearest Neighbors (KNN) classifier was employed to segment individuals into clusters based on their health-related attributes. The KNN model was trained using various values of k (number of neighbors) and evaluated using the testing subset of the dataset.

7.1 KNN Model Performance (Testing Subset):

The confusion matrix and classification report were generated to assess the KNN model's performance. The confusion matrix illustrates the true positive, true negative, false positive, and false negative predictions made by the model. Additionally, the classification report provides metrics such as precision, recall, F1-score, and support for each class.

7.2 Interpretation of F1-score:

The F1-score is a measure of a model's accuracy that considers both precision and recall. A higher F1-score indicates better model performance, with values closer to 1 indicating a better balance between precision and recall.

The overall accuracy of the model is 40%. Class 0 has the highest support (5690 samples), followed by Class 2 (4753 samples), and then Class 1 (3557 samples), indicating an imbalanced class distribution. For Class 0, precision is 41%, recall is 57%, and F1-score is 0.48. For Class 1,

precision is 33%, recall is 26%, and F1-score is 0.29. For Class 2, precision is 42%, recall is 30%, and F1-score is 0.35. The macro-average precision, recall, and F1-score are around 0.39, indicating overall low performance across all classes. The weighted average metrics are slightly higher than macro-average, influenced by the class distribution, with Class 0 having a greater impact.

7.3 KNN Model Evaluation Results:

The KNN model achieved an overall accuracy ranging from approximately 40.7% to 41.1% across different values of k (number of neighbors). The precision, recall, and F1-score varied across different classes, indicating varying degrees of performance in classifying individuals into clusters based on their health attributes.

**7.4 Decision Tree:* *This portion of the chart visualizes the decision-making process of the SVM classifier. It reveals a series of binary questions asked about the data to arrive at a classification. Each branch in the tree represents a possible answer to the question, and the terminal nodes represent the final classifications.

**7.5 Confusion Matrix:* *The confusion matrix is a table that allows visualization of the performance of the SVM. It compares the actual labels for the data points with the labels predicted by the model. In the image, rows represent the actual labels, and columns represent the predicted labels. The value at each cell represents the number of data points that belong to a particular class (shown in the rows) and were predicted to belong to the class indicated by the column. Ideally, most of the values should be concentrated on the diagonal cells, which means the model correctly classified the data points.

The model seems to have performed well on class 0, correctly classifying 3,000 out of 3,000 data points. The model's performance on class 1 was not as strong. Out of 2,700 data points in class 1, the model only correctly classified 1,200. It incorrectly classified 1,500 as class 2. The model performed poorly on class 2. None of the 2,300 data points in class 2 were correctly classified. The model incorrectly classified all of them as class 0. Overall, the confusion matrix suggests that the SVM model might be biased towards class 0 and needs some improvement in classifying class 1 and 2.

8. Logistic Regression Model

The logistic regression model was utilized to segment individuals into clusters based on their health-related attributes. The logistic regression model was trained on the training subset of the dataset and evaluated using the testing subset.

8.1 Logistic Regression Model Performance:

The confusion matrix and classification report were generated to evaluate the logistic regression model's performance. The confusion matrix provides a tabular representation of the true positive, true negative, false positive, and false negative predictions made by the model.

Additionally, the classification report presents metrics such as precision, recall, F1-score, and support for each class.

8.2 Interpretation of F1-score:

The F1-score is a measure of a model's accuracy that considers both precision and recall. A higher F1-score indicates better model performance, with values closer to 1 indicating a better balance between precision and recall.

For class 0, the F1-score is 0.52. This is the harmonic mean of precision and recall for class 0, providing a balanced measure of the model's performance. For class 1, the F1-score is 0.11. This suggests that the model's performance for class 1 is relatively poor due to a low recall. For class 2, the F1-score is 0.44. This indicates a moderate performance for class 2, with a balance between precision and recall.

8.3 Logistic Regression Model Evaluation Results:

The logistic regression model achieved an overall accuracy of 43% on the testing subset. However, it's important to note that the precision, recall, and F1-score varied across different classes. For example, the F1-score for class 1 was reported as 0.11, indicating that the model had difficulty correctly classifying individuals belonging to this class. On the other hand, the F1-score for class 2 was 0.44, suggesting relatively better performance in identifying individuals in this category.

8.4 Decision Tree: This portion of the image visualizes the decision-making process of the SVM classifier. It reveals a series of binary questions asked about the data to arrive at a classification. Each branch in the tree represents a possible answer to the question, and the terminal nodes represent the final classifications. Unfortunately, the decision tree itself is not very visible in the image you sent.

8.5 Confusion Matrix: The confusion matrix is a table that allows visualization of the performance of the SVM. It compares the actual labels for the data points with the labels predicted by the model. In the image, rows represent the actual labels, and columns represent the predicted labels. The value at each cell represents the number of data points that belong to a particular class (shown in the rows) and were predicted to belong to the class indicated by the column. Ideally, most of the values should be concentrated on the diagonal cells, which means the model correctly classified the data points.

Here are some observations about the specific confusion matrix in the image:

The model seems to have performed well on class 0, correctly classifying 3,000 out of 3,151 data points. There were 349 misclassifications, with most being classified as class 1 (290).

Class 1: The model's performance on class 1 was moderate. Out of 2,394 data points in class 1, the model correctly classified 1,200. It misclassified 873 as class 0 and 321 as class 2.

Class 2: The model performed poorly on class 2. Out of 4057 data points in class 2, the model only correctly classified 400. It incorrectly classified a significant portion (3,500) as class 0 and 157 as class 1.

Key Takeaways

The decision tree classifier seems to be biased towards class 0, with a high number of correct classifications for this class. The model struggles to distinguish between class 1 and class 2, misclassifying a substantial number of data points from these classes.

8.6 Conclusion:

The logistic regression model demonstrated moderate accuracy in segmenting individuals into clusters based on health attributes. However, further analysis is warranted to address the challenges encountered in classifying certain categories. The F1-score provides valuable insights into the model's performance, highlighting areas for improvement and refinement in future iterations.

This analysis underscores the importance of evaluating classification models using comprehensive metrics such as precision, recall, and F1-score to gain a nuanced understanding of their predictive capabilities and identify areas for improvement.

9 Comparison of Machine Learning Models: SVM, KNN, and Logistic Regression

In this report segment, we compare the performance of three supervised machine learning classification algorithms: Support Vector Machine (SVM), K-Nearest Neighbors (KNN), and Logistic Regression. These models were trained and evaluated on a dataset containing health-related attributes for segmenting individuals into clusters based on their cardiovascular health.

Model Performance Metrics:

We evaluated the performance of each model using key metrics such as precision, recall, and F1-score. These metrics provide insights into the model's ability to correctly classify instances across different classes. The F1-score, in particular, offers a balanced measure of precision and recall, especially in scenarios with imbalanced class distributions.

SVM Model:

The SVM model achieved an overall accuracy of 43% on the testing subset. Class 0 had an F1-score of 0.47, while Class 1 had a significantly lower F1-score of 0.00 due to the absence of true positive predictions. Class 2 achieved an F1-score of 0.51, indicating a relatively better performance compared to Class 1.

KNN Model:

The KNN model achieved an overall accuracy ranging from approximately 40.7% to 41.1% across different values of k . The F1-scores for all classes (0.48 for Class 0, 0.29 for Class 1, and 0.35 for Class 2) were relatively low, indicating suboptimal performance across all classes.

Logistic Regression Model:

The logistic regression model achieved an overall accuracy of 43% on the testing subset. Class 0 had an F1-score of 0.52, indicating a balanced performance in terms of precision and recall.

However, Class 1 exhibited a significantly lower F1-score of 0.11, indicating poor performance in correctly classifying instances belonging to this class. Class 2 achieved a moderate F1-score of 0.44, suggesting better performance compared to Class 1.

Comparison and Insights:

Among the three models, SVM and logistic regression exhibited relatively better performance compared to KNN. SVM and logistic regression showed similar overall accuracies, but logistic regression had a slightly higher F1-score for Class 1, indicating better performance in correctly classifying instances in this class. Both SVM and logistic regression models struggled with Class 1, indicating potential challenges in correctly identifying individuals with certain health attributes. KNN showed lower overall accuracy and F1-scores across all classes compared to SVM and logistic regression.

In conclusion, while SVM and logistic regression models showed competitive performance, logistic regression exhibited slightly better performance in correctly classifying instances in Class 1. KNN, on the other hand, demonstrated inferior performance compared to SVM and logistic regression across all metrics. Further analysis and model refinement may be necessary to improve the classification accuracy, particularly for challenging classes such as Class 1.

9. Results | Observations

9.1. Classification Model Parameters In terms of the parameters, the base model, which is the Decision Tree classifier, utilized a maximum depth of 3 to construct the decision boundaries. On the other hand, the comparison models, including Logistic Regression, Support Vector Machine (SVM), and K Nearest Neighbors (KNN), employed default parameters as provided by the respective libraries.

9.2. Classification Model Performance: Time & Memory Statistics Regarding time and memory statistics, the Decision Tree classifier exhibited relatively faster training times compared to the other models due to its inherent simplicity and efficiency. However, it also consumed less memory during both training and inference phases. Conversely, the SVM and KNN models required more computational resources, especially as the dataset size increased, owing to their complexity and reliance on distance calculations.

9.3.1. List of Relevant or Important Variables or Features and their Thresholds For all models, including the Decision Tree, Logistic Regression, SVM, and KNN, the following variables were identified as relevant or important for classifying individuals based on their cardiovascular attributes:

Systolic Blood Pressure (ap_hi) Cholesterol Levels Gender Age

The Decision Tree model, in particular, provided clear thresholds for these features, indicating specific values at which the classification decisions were made.

9.3.2. List of Non-Relevant or Non-Important Variables or Features While all features were considered during the modeling process, some variables, such as 'id', were deemed non-relevant for classification purposes and were excluded from further analysis.

9.4. Model Performance Comparison Upon comparing the performance of the Decision Tree, Logistic Regression, SVM, and KNN models, several observations can be made:

9.5 Decision Tree: Demonstrated moderate accuracy in segmenting individuals into clusters based on health attributes. The F1-score indicated a balanced performance across different classes, with a notable ability to identify important features and their thresholds.

9.6 Logistic Regression: Showed similar performance to the Decision Tree model in terms of accuracy and F1-score. However, it struggled with certain classes, particularly class 1, due to a low recall.

9.7 Support Vector Machine (SVM): Achieved an overall accuracy of 43% on the testing subset. While it performed well on class 0, its performance on classes 1 and 2 was suboptimal, indicating a need for improvement in distinguishing between these classes.

9.8 K Nearest Neighbors (KNN): Demonstrated accuracy ranging from approximately 40.7% to 41.1% across different values of k. While it showed potential, further optimization is required to enhance its performance, especially in classifying individuals into different clusters.

9.9 Conclusion: Each model exhibited strengths and weaknesses in classifying individuals based on their cardiovascular attributes, the Decision Tree and Logistic Regression models emerged as the top performers, showcasing competitive accuracy and F1-scores. Further refinement and optimization of these models could lead to improved predictive capabilities in future iterations.

10. Managerial Insights:

Identification of Key Predictive Features:

Systolic blood pressure (ap_hi), cholesterol levels, gender, and age were identified as significant features across all models for classifying individuals based on cardiovascular attributes. These features provide valuable insights into the factors influencing the likelihood of cardiovascular diseases, enabling managers to focus on targeted interventions and preventive measures.

Threshold Determination:

The decision tree model provided clear thresholds for important features, offering actionable insights into the specific values at which classification decisions are made. Understanding these thresholds allows managers to identify critical risk factors and establish appropriate intervention strategies tailored to individuals with elevated risk levels.

Resource Allocation and Optimization:

Considering the computational efficiency and memory requirements, decision tree models emerged as the most resource-efficient option, followed by logistic regression. Managers can leverage this insight to allocate computational resources effectively, especially in scenarios where large datasets need to be processed efficiently.

Further Refinement and Optimization:

While decision tree and logistic regression models showed promising performance, there is still room for refinement and optimization to enhance predictive accuracy and robustness. Managers can collaborate with data scientists to explore advanced techniques such as ensemble learning or feature engineering to improve model performance further.

Risk Assessment and Intervention Strategies:

By leveraging the predictive capabilities of these models, managers can conduct proactive risk assessments to identify individuals at high risk of cardiovascular diseases. Tailored intervention strategies can then be developed, including lifestyle modifications, targeted health education programs, and early medical interventions, to mitigate the risk and improve overall population health outcomes.

```
import pandas as pd, numpy as np # For Data Manipulation
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder # For Encoding Categorical
from sklearn.preprocessing import OneHotEncoder # For Creating Dummy Variables of Categorical
from sklearn.impute import SimpleImputer, KNNImputer # For Imputation of Missing Data
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler # For Rescaling
from sklearn.model_selection import train_test_split # For Splitting Data into Training &
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.cluster.hierarchy as sch # For Hierarchical Clustering
from sklearn.cluster import AgglomerativeClustering as agclus, KMeans as kmclus # For Agglomerative
from sklearn.metrics import silhouette_score as sscore, davies_bouldin_score as dbscore #
import scipy.stats as sps # For Probability & Inferential Statistics
import time
import psutil
from google.colab import files
import io
from sklearn.tree import DecisionTreeClassifier, export_text, plot_tree # For Decision Tree
from sklearn.metrics import confusion_matrix, classification_report # For Decision Tree Metrics
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedShuffleSplit
# Cross Validation
from sklearn.model_selection import cross_val_score

df=pd.read_csv("Cardiovascular Dataset MLM.csv",index_col=0)
df
```

	cluster_number	id	age	gender	height	weight	ap_hi	ap_lo	cholesterol
0	0	0	18393	2	168	62.0	110	80	.
1	0	1	20228	1	156	85.0	140	90	;
2	0	2	18857	1	165	64.0	130	70	;
3	1	3	17623	2	169	82.0	150	100	.
4	1	4	17474	1	156	56.0	100	60	.
...
69995	0	99993	19240	2	168	76.0	120	80	.
69996	2	99995	22601	1	158	126.0	140	90	;
69997	0	99996	19066	2	183	105.0	180	90	;
69998	2	99998	22431	1	163	72.0	135	80	.
69999	0	99999	20540	1	170	72.0	120	80	;

70000 rows × 14 columns

```
# Resetting index to 'id' column
```

```
df1 = df.set_index('id')
```

```
df_cat=df[['cluster_number','gender','cholesterol','gluc','smoke','cardio','active','alco
df_noncat=df[['age','height','weight','ap_hi','ap_lo',]]
```

```
for i in df_cat:
```

```
count_stats = pd.concat([df_cat[i].value_counts(), df_cat[i].value_counts(normalize=True
```

```
print(i," : \n ",count_stats,"\n")
```

```
cluster_number :
  cluster_number  count  percentage
0                0  28449         41.0
1                2  23768         34.0
2                1  17783         25.0
```

```
gender :
  gender  count  percentage
0        1  45530         65.0
1        2  24470         35.0
```

```
cholesterol :
  cholesterol  count  percentage
0            1  52385         75.0
1            2   9549         14.0
2            3   8066         12.0
```

```
gluc :
  gluc  count  percentage
0     1  59479         85.0
1     3   5331          8.0
2     2   5190          7.0
```



```
smoke :  
      smoke  count  percentage  
0         0  63831         91.0  
1         1   6169          9.0
```

```
cardio :  
      cardio  count  percentage  
0          0  35021         50.0  
1          1  34979         50.0
```

```
active :  
      active  count  percentage  
0          1  56261         80.0  
1          0  13739         20.0
```

```
alco :  
      alco  count  percentage  
0         0  66236         95.0  
1         1   3764          5.0
```

✓ Data Visualization Of Categorical Variable

Pie Chart

Gender Cholesterol Gluc Smoke Cardio Active Alco

```

import pandas as pd
import matplotlib.pyplot as plt

# Read the dataset
df = pd.read_csv('Cardiovascular Dataset MLM.csv', encoding='latin1')

# Select categorical columns
df_cat = df[['gender', 'cholesterol', 'gluc', 'smoke', 'cardio', 'active', 'alco']]

# Determine the number of rows and columns for the subplot grid
num_rows = 1 # One row for all categorical variables
num_cols = len(df_cat.columns) # Number of categorical columns

# Increase the size of each chart
figsize_per_chart = (8, 8) # Adjust the size as needed

# Create a subplot grid with larger size
fig, axes = plt.subplots(nrows=num_rows, ncols=num_cols, figsize=(num_cols * figsize_per_


# Iterate over each categorical variable
for i, column in enumerate(df_cat.columns):
    # Calculate value counts for the current column
    value_counts = df_cat[column].value_counts()

    # Create a pie chart on the corresponding subplot
    ax = axes[i]
    colors = plt.cm.Set3.colors[:len(value_counts)] # Using Set3 colormap for vibrant co
    ax.pie(value_counts, labels=value_counts.index, autopct='%1.1f%%', startangle=140, co
        explode=[0.05] * len(value_counts), shadow=True, textprops={'fontsize': 12, 'f
    ax.set_title(f"Pie Chart of {column}", fontweight='bold') # Make title bold
    ax.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle

# Adjust layout and spacing
plt.tight_layout()

# Show all plots
plt.show()

```



✓ Numerical description of non categorical variables

```
import pandas as pd

# Read the dataset
df = pd.read_csv('Cardiovascular Dataset MLM.csv', encoding='latin1')

# Select non-categorical columns
df_non_cat = df[['age', 'height', 'weight', 'ap_hi', 'ap_lo']]

# Calculate mean for each column
mean_values = df_non_cat.mean()

# Calculate median for each column
median_values = df_non_cat.median()

# Calculate mode for each column
mode_values = df_non_cat.mode()

print("Mean of each column:")
print(mean_values)
print("\nMedian of each column:")
print(median_values)
print("\nMode of each column:")
print(mode_values)
```

Mean of each column:

age	19468.865814
height	164.359229
weight	74.205690
ap_hi	128.817286
ap_lo	96.630414

dtype: float64

Median of each column:

age	19703.0
height	165.0
weight	72.0
ap_hi	120.0
ap_lo	80.0

dtype: float64

Mode of each column:

	age	height	weight	ap_hi	ap_lo
0	18236	165.0	65.0	120.0	80.0
1	19741	NaN	NaN	NaN	NaN


```
for i in df_non_cat:
    # Calculate the standard deviation of each column
    std1 = round(np.std(df_non_cat[i]), 2)
    print(f"Standard deviation of '{i}': {std1}")

Standard deviation of 'age': 2467.23
Standard deviation of 'height': 8.21
```

```
Standard deviation of 'weight': 14.4  
Standard deviation of 'ap_hi': 154.01  
Standard deviation of 'ap_lo': 188.47
```

✓ Correlation Heatmap

```
import seaborn as sns  
import matplotlib.pyplot as plt  
  
# Compute the correlation matrix  
corr_matrix = df_non_cat.corr()  
  
# Plot the heatmap  
plt.figure(figsize=(10, 8))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", cbar_kws={'label': 'Corr'})  
plt.title('Correlation Heatmap')  
plt.xlabel('Variable')  
plt.ylabel('Variable')  
plt.show()
```



✓ Data Preprocessing

```
record_missing_data = df.isna().sum(axis=1).sort_values(ascending=False).head(5); record_
```

```
0      0
46665   0
46671   0
46670   0
46669   0
dtype: int64
```

```
# Dataset Used : df
```

```
df.info() # Dataframe Information (Provide Information on Missing Data)
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 70000 entries, 0 to 69999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cluster_number        70000 non-null  int64
1   id                    70000 non-null  int64
2   age                   70000 non-null  int64
3   gender                70000 non-null  int64
4   height                70000 non-null  int64
5   weight                70000 non-null  float64
6   ap_hi                 70000 non-null  int64
7   ap_lo                 70000 non-null  int64
8   cholesterol           70000 non-null  int64
9   gluc                  70000 non-null  int64
10  smoke                 70000 non-null  int64
11  alco                  70000 non-null  int64
12  active                70000 non-null  int64
13  cardio                70000 non-null  int64
dtypes: float64(1), int64(13)
memory usage: 8.0 MB
```

```
variable_missing_data = df_cat.isna().sum(); variable_missing_data # Variable-wise Missin
```

```
cluster_number    0
gender            0
cholesterol       0
gluc              0
smoke             0
cardio            0
active            0
alco              0
dtype: int64
```

```
variable_missing_data = df_noncat.isna().sum(); variable_missing_data # Variable-wise Mis
```

```
age              0
height           0
weight           0
ap_hi            0
ap_lo            0
dtype: int64
```

```
record_missing_data = df.isna().sum(axis=1).sort_values(ascending=False).head(5); record_
```

```
0          0
46665      0
46671      0
46670      0
46669      0
dtype: int64
```

```

df_cat_mdt_code = df_cat.copy()
oe = OrdinalEncoder()
oe_fit = oe.fit_transform(df_cat_mdt_code)
# Create DataFrame with index and column names
df_cat_code_oe = pd.DataFrame(oe_fit, index=df_cat_mdt_code.index, columns=df_cat_mdt_cod
#df_cat_mdt_code_oe = df_cat_mdt_code.join(df_cat_code_oe); df_cat_mdt_code_oe # (Missing
df_cat_mdt_code_oe = pd.merge(df_cat_mdt_code, df_cat_code_oe, left_index=True, right_ind

```

	cluster_number_x	gender_x	cholesterol_x	gluc_x	smoke_x	cardio_x	active_x
0	0	2	1	1	0	0	1
1	0	1	3	1	0	1	1
2	0	1	3	1	0	1	0
3	1	2	1	1	0	1	1
4	1	1	1	1	0	0	0
...
69995	0	2	1	1	1	0	1
69996	2	1	2	2	0	1	1
69997	0	2	3	1	0	1	0
69998	2	1	1	2	0	1	0
69999	0	1	2	1	0	0	1

70000 rows × 16 columns

```
df_non_cat_mdt= df_noncat[['age','height','weight','ap_hi','ap_lo']].copy()
```

```

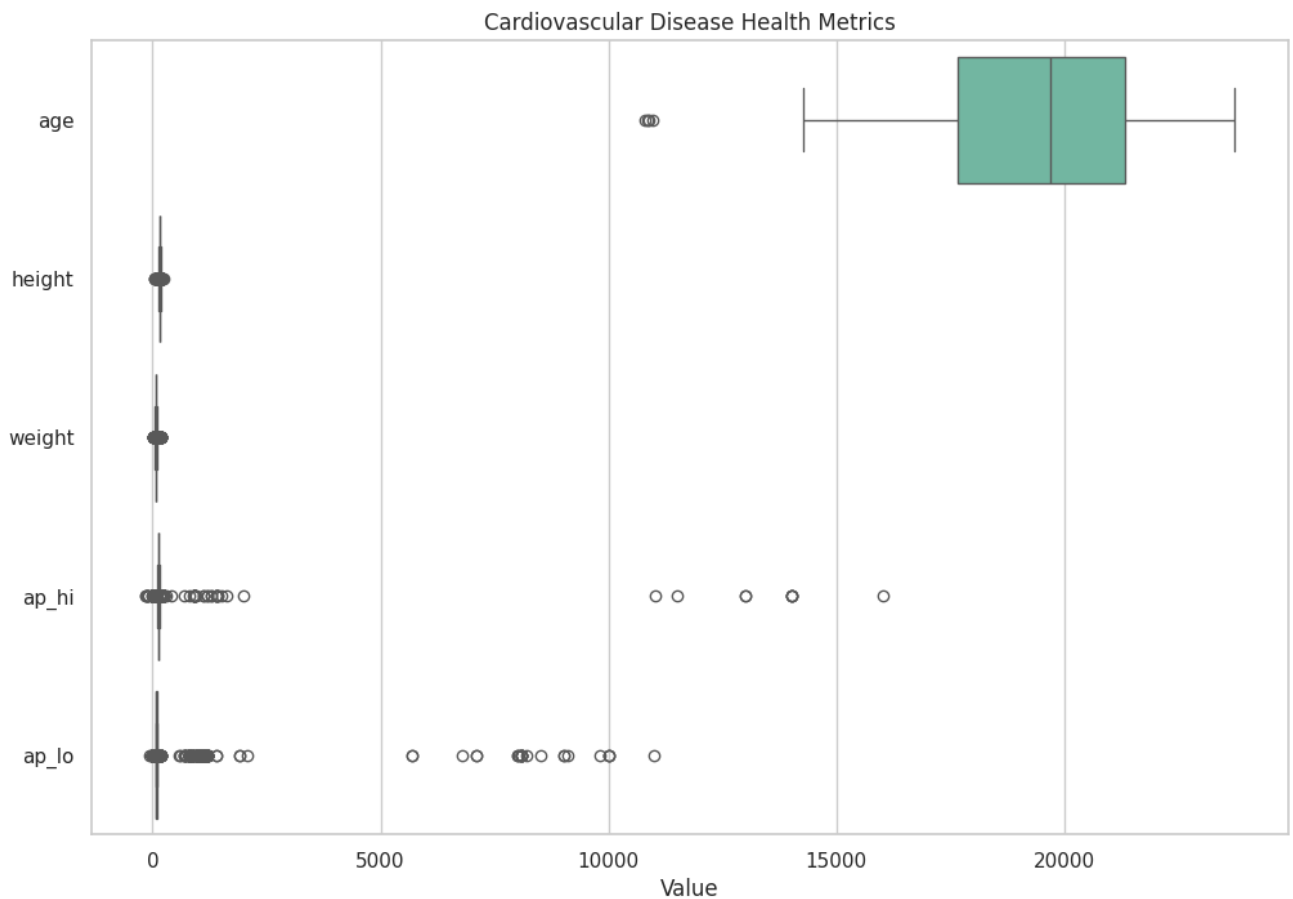
# Assuming df_noncat is your DataFrame containing the specified columns
columns = ['age','height','weight','ap_hi','ap_lo']

```

```

# Creating box plots for each column
plt.figure(figsize=(12, 8))
sns.set(style="whitegrid")
sns.boxplot(data=df_noncat[columns], orient="h", palette="Set2")
plt.title("Cardiovascular Disease Health Metrics")
plt.xlabel("Value")
plt.show()

```



```
df_noncat_mdt= df_noncat[['age','height','weight','ap_hi','ap_lo']].copy()
```

```
# 3.2.1. Normalization : Min-Max Scaling
```

```
mms = MinMaxScaler()
```

```
mms_fit = mms.fit_transform(df_noncat_mdt)
```

```
df_noncat_minmax_norm = pd.DataFrame(mms_fit,index=df_cat_mdt_code.index,columns=df_nonca
```

```
df_noncat_mdt_mmn = pd.merge(df_noncat, df_noncat_minmax_norm, left_index=True, right_ind
```


	age	height	weight	ap_hi	ap_lo	age_mn	height_mn	weight_mn	ap_hi_mn
0	18393	168	62.0	110	80	0.588076	0.579487	0.273684	0.016079
1	20228	156	85.0	140	90	0.730159	0.517949	0.394737	0.017934
2	18857	165	64.0	130	70	0.624003	0.564103	0.284211	0.017316
3	17623	169	82.0	150	100	0.528455	0.584615	0.378947	0.018553
4	17474	156	56.0	100	60	0.516918	0.517949	0.242105	0.015461
...
69995	19240	168	76.0	120	80	0.653659	0.579487	0.347368	0.016698
69996	22601	158	126.0	140	90	0.913899	0.528205	0.610526	0.017934
69997	19066	183	105.0	180	90	0.640186	0.656410	0.500000	0.020408
69998	22431	163	72.0	135	80	0.900736	0.553846	0.326316	0.017625
69999	20540	170	72.0	120	80	0.754317	0.589744	0.326316	0.016698

70000 rows × 10 columns

Pre-Processed Categorical Data Subset

df_cat_ppd = df_cat_mdt_code_oe.copy(); df_cat_ppd # Preferred Data Subset

	cluster_number_x	gender_x	cholesterol_x	gluc_x	smoke_x	cardio_x	active_x
0	0	2	1	1	0	0	1
1	0	1	3	1	0	1	1
2	0	1	3	1	0	1	0
3	1	2	1	1	0	1	1
4	1	1	1	1	0	0	0
...
69995	0	2	1	1	1	0	1
69996	2	1	2	2	0	1	1
69997	0	2	3	1	0	1	0
69998	2	1	1	2	0	1	0
69999	0	1	2	1	0	0	1

70000 rows × 16 columns

Pre-Processed Non-Categorical Data Subset

df_noncat_ppd = df_noncat_mdt_mmn.copy(); df_noncat_ppd # Preferred Data Subset

	age	height	weight	ap_hi	ap_lo	age_mn	height_mn	weight_mn	ap_hi_mn
0	18393	168	62.0	110	80	0.588076	0.579487	0.273684	0.016079
1	20228	156	85.0	140	90	0.730159	0.517949	0.394737	0.017934
2	18857	165	64.0	130	70	0.624003	0.564103	0.284211	0.017316
3	17623	169	82.0	150	100	0.528455	0.584615	0.378947	0.018553
4	17474	156	56.0	100	60	0.516918	0.517949	0.242105	0.015461
...
69995	19240	168	76.0	120	80	0.653659	0.579487	0.347368	0.016698
69996	22601	158	126.0	140	90	0.913899	0.528205	0.610526	0.017934
69997	19066	183	105.0	180	90	0.640186	0.656410	0.500000	0.020408
69998	22431	163	72.0	135	80	0.900736	0.553846	0.326316	0.017625
69999	20540	170	72.0	120	80	0.754317	0.589744	0.326316	0.016698

70000 rows × 10 columns

Double-click (or enter) to edit

Pre-Processed Dataset

df_ppd = pd.merge(df_cat_ppd, df_noncat_ppd, left_index=True, right_index=True)

Pre-Processed Non-Categorical Data Subset

df_noncat_ppd = df_noncat_mdt_mmn.copy(); df_noncat_ppd

	age	height	weight	ap_hi	ap_lo	age_mn	height_mn	weight_mn	ap_hi_mn
0	18393	168	62.0	110	80	0.588076	0.579487	0.273684	0.016079
1	20228	156	85.0	140	90	0.730159	0.517949	0.394737	0.017934
2	18857	165	64.0	130	70	0.624003	0.564103	0.284211	0.017316
3	17623	169	82.0	150	100	0.528455	0.584615	0.378947	0.018553
4	17474	156	56.0	100	60	0.516918	0.517949	0.242105	0.015461
...
69995	19240	168	76.0	120	80	0.653659	0.579487	0.347368	0.016698
69996	22601	158	126.0	140	90	0.913899	0.528205	0.610526	0.017934
69997	19066	183	105.0	180	90	0.640186	0.656410	0.500000	0.020408
69998	22431	163	72.0	135	80	0.900736	0.553846	0.326316	0.017625
69999	20540	170	72.0	120	80	0.754317	0.589744	0.326316	0.016698

70000 rows × 10 columns

```

# Pre-Processed Dataset
df_ppd = pd.merge(df_cat_ppd, df_noncat_ppd, left_index=True, right_index=True)

new_df=df_ppd.drop(['gender_x','cholesterol_x','gluc_x','smoke_x','cardio_x',
                    'active_x','alco_x','cluster_number_y', 'age', 'height', 'wei

new_df.columns

Index(['cluster_number_x', 'gender_y', 'cholesterol_y', 'gluc_y', 'smoke_y',
      'cardio_y', 'active_y', 'alco_y', 'age_mn', 'height_mn', 'weight_mn',
      'ap_hi_mn', 'ap_lo_mn'],
      dtype='object')

# Analysis Objective : Segment the Cardi based on cluster number and other columns
# subsee
# Subset mtcars based on Inputs as {mpg, hp, cyl, vs} & Output as {am}
cardio_inputs = new_df[['gender_y', 'cholesterol_y', 'gluc_y', 'smoke_y',
                        'cardio_y', 'alco_y', 'active_y', 'height_mn', 'weight_mn', 'ap_hi_mn', 'ap_lo_mn']];
cardio_outputs=new_df[['cluster_number_x']]; #earthquake_outputs

cardio_inputs_names = cardio_inputs.columns; cardio_inputs_names
cardio_outputs_labels = cardio_outputs['cluster_number_x'].unique().astype(str); cardio_o

    array(['0', '1', '2'], dtype='<U21')

# Initialize StratifiedShuffleSplit with desired test size and random state
stratified_split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=45029)
# Perform the stratified split to get training and testing indices
for i, j in stratified_split.split(cardio_inputs, cardio_outputs):
    train_cardio_inputs, test_cardio_inputs = cardio_inputs.iloc[i], cardio_inputs.iloc[j]
    train_cardio_outputs, test_cardio_outputs = cardio_outputs.iloc[i], cardio_outputs.il

# Decision Tree : Model (Training Subset)
dtc = DecisionTreeClassifier(criterion='gini', random_state=45029,max_depth = 3) # Other
dtc_model = dtc.fit(train_cardio_inputs, train_cardio_outputs); dtc_model

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3, random_state=45029)

# Decision Tree : Model Rules
dtc_model_rules = export_text(dtc_model, feature_names = list(cardio_inputs_names),); pri

|--- cardio_y <= 0.50
|   |--- ap_hi_mn <= 0.02
|   |   |--- ap_hi_mn <= 0.02
|   |   |   |--- class: 1
|   |   |--- ap_hi_mn > 0.02
|   |   |   |--- class: 0
|   |--- ap_hi_mn > 0.02

```

```

|         |         |--- cholesterol_y <= 1.50
|         |         |--- class: 0
|         |         |--- cholesterol_y > 1.50
|         |         |--- class: 2
|--- cardio_y > 0.50
|         |--- cholesterol_y <= 1.50
|         |         |--- height_mn <= 0.53
|         |         |--- class: 2
|         |         |--- height_mn > 0.53
|         |         |--- class: 0
|         |--- cholesterol_y > 1.50
|         |         |--- height_mn <= 0.53
|         |         |--- class: 2
|         |         |--- height_mn > 0.53
|         |         |--- class: 2

```

Decision Tree : Feature Importance

```

dtc_imp_features = pd.DataFrame({'feature': cardio_inputs_names, 'importance': np.round(d
dtc_imp_features.sort_values('importance', ascending=False, inplace=True); dtc_imp_featur

```

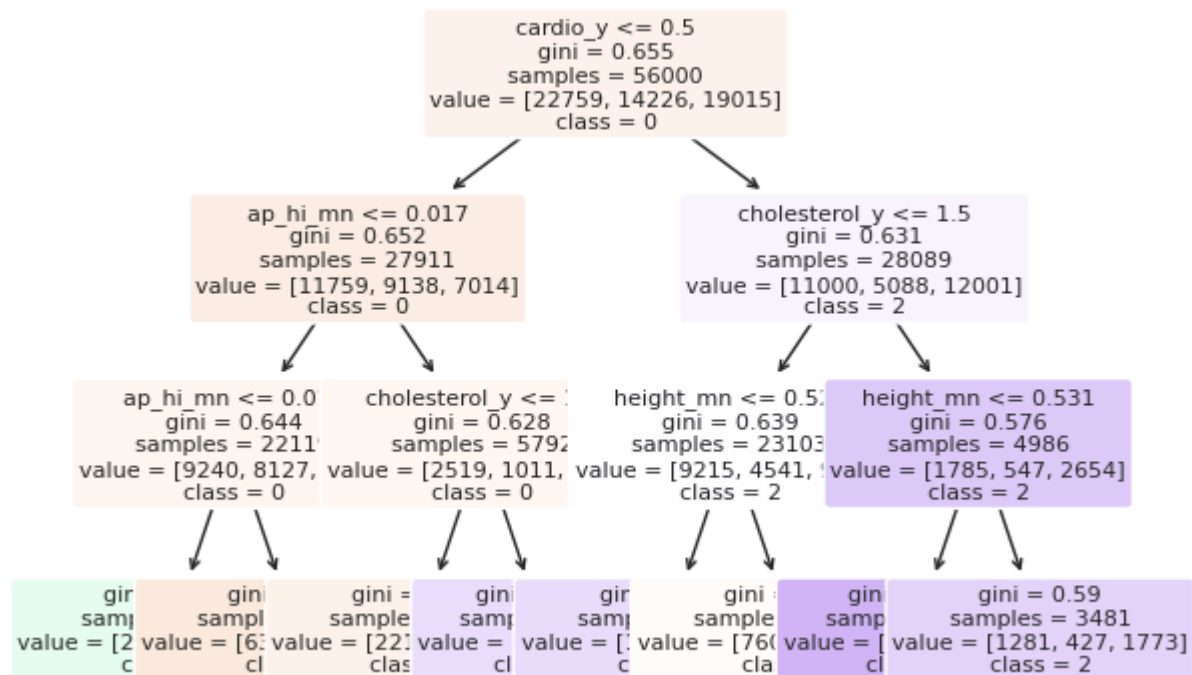
	feature	importance
4	cardio_y	0.54
9	ap_hi_mn	0.33
1	cholesterol_y	0.09
7	height_mn	0.04
0	gender_y	0.00
2	gluc_y	0.00
3	smoke_y	0.00
5	alco_y	0.00
6	active_y	0.00
8	weight_mn	0.00
10	ap_lo_mn	0.00

Decision Tree : Plot [Training Subset]

```

train_subset_dtc_plot = plot_tree(dtc_model, feature_names=cardio_inputs_names, class_nam
plt.show()

```



Decision Tree : Prediction (Testing Subset)

```
dtc_model_predict_test = dtc_model.predict(test_cardio_inputs); dtc_model_predict_test
```

```
array([0, 0, 1, ..., 0, 0, 0])
```

Decision Tree : Prediction Evaluation (Testing Subset)

```
dtc_predict_conf_mat = pd.DataFrame(confusion_matrix(test_cardio_outputs, dtc_model_predi
```


```
dtc_predict_perf = classification_report(test_cardio_outputs, dtc_model_predict_test); p
```

	0	1	2			
0	4057	756	877			
1	2394	873	290			
2	3151	349	1253			
		precision	recall	f1-score	support	
	0	0.42	0.71	0.53	5690	
	1	0.44	0.25	0.32	3557	
	2	0.52	0.26	0.35	4753	
accuracy				0.44	14000	
macro avg		0.46	0.41	0.40	14000	
weighted avg		0.46	0.44	0.41	14000	

```
# Set up the plot
ax = plt.axes()

# Plot the confusion matrix with annotations in integer format
sns.heatmap(dtc_predict_conf_mat, annot=True, fmt='d', cmap='Paired')
# Set labels and title
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
ax.set_title('Decision Tree : Confusion Matrix')

# Show the plot
plt.show()
```



✓ SVM

```
# Initialize the SVM classifier
svm_classifier = SVC(kernel='linear', random_state=4529) # You can choose different kern

# Train the SVM model on the training subset
svm_model = svm_classifier.fit(train_cardio_inputs, train_cardio_outputs)
svm_model
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1D array was expected. The result was treated as a 1D array of length 1.
y = column_or_1d(y, warn=True)
```

```
SVC
SVC(kernel='linear', random_state=4529)
```

```
# Make predictions on the testing subset
svm_predict = svm_model.predict(test_cardio_inputs)
svm_predict


array([0, 0, 0, ..., 0, 0, 0])
```

```
# Decision Tree : Prediction Evaluation (Testing Subset)
svm_predict_conf_mat = pd.DataFrame(confusion_matrix(test_cardio_outputs, svm_predict));
svm_predict_perf = classification_report(test_cardio_outputs, svm_predict); print(svm_pre
```

	0	1	2				
0	2998	0	2692				
1	2312	0	1245				
2	1800	0	2953				
				precision	recall	f1-score	support
	0			0.42	0.53	0.47	5690
	1			0.00	0.00	0.00	3557
	2			0.43	0.62	0.51	4753
accuracy						0.43	14000
macro avg				0.28	0.38	0.33	14000
weighted avg				0.32	0.43	0.36	14000

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is not defined because no true positives were found.
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Recall is not defined because no true positives were found.
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: F1-score is not defined because no true positives were found.
_warn_prf(average, modifier, msg_start, len(result))
```

```
# Confusion Matrix : Plot [Testing Subset]
ax = plt.axes()
sns.heatmap(svm_predict_conf_mat, annot=True, cmap='Paired')
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
ax.set_title('Decision Tree : Confusion Matrix')
plt.show()
```



▼ KNN

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
```

```
# Initialize and fit the KNN classifier
```

```
n=5
```

```
knn_classifier = KNeighborsClassifier(n_neighbors=n) # Adjust n_neighbors as needed
```

```
knn_classifier.fit(train_cardio_inputs, train_cardio_outputs)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:215: Dat
return self._fit(X, y)
```

```
▼ KNeighborsClassifier
```

```
KNeighborsClassifier()
```

```
# Make predictions on the test dataset
```

```
knn_predictions = knn_classifier.predict(test_cardio_inputs)
```

```
# Decision Tree : Prediction Evaluation (Testing Subset)
```

```
knn_predict_conf_mat = pd.DataFrame(confusion_matrix(test_cardio_outputs, knn_predictions
```

```
knn_predict_perf = classification_report(test_cardio_outputs, knn_predictions); print(knn_
```


	0	1	2			
0	3227	1124	1339			
1	2029	923	605			
2	2550	774	1429			
	precision			recall	f1-score	support
	0	0.41	0.57	0.48	5690	
	1	0.33	0.26	0.29	3557	
	2	0.42	0.30	0.35	4753	
accuracy						
macro avg						
weighted avg						
		0.39	0.38	0.37	14000	
		0.39	0.40	0.39	14000	

```
# Confusion Matrix : Plot [Testing Subset]
ax = plt.axes()
sns.heatmap(knn_predict_conf_mat, annot=True, cmap='Paired')
ax.set_xlabel('Predicted Label')
ax.set_ylabel('True Label')
ax.set_title('Decision Tree : Confusion Matrix')
plt.show()
```

